

CS3802--Machine Learning Algorithms Lab

Adithya V | BTech CSE (IoT) - A | 21011102009

Exercise 6

Use the teleco-customer-churn dataset for the following:

1. Use the attached file and run SVM, Decision tree, Random Forest and any one boosting algorithm.
2. Find out the different tunable parameters for each algorithms mentioned above.
3. Apply gridsearchCV and randomizedsearchCV for all the above classification algorithms and get the best parameters.

Importing necessary libraries and reading the dataset

```
In [ ]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import RandomizedSearchCV
```

```
c:\Python311\Lib\site-packages\numpy\_distributor_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
c:\Python311\Lib\site-packages\numpy\.libs\libopenblas64__v0.3.21-gcc_10_3_0.dll
c:\Python311\Lib\site-packages\numpy\.libs\libopenblas64__v0.3.23-gcc_10_3_0.dll
warnings.warn("loaded more than 1 DLL from .libs:")
```

```
In [ ]: df = pd.read_csv("Telco-Customer-Churn.csv")
df = df.drop('customerID', axis = 1)
df.head()
```

```
Out[ ]:
```

| | gender | SeniorCitizen | Partner | Dependents | tenure | PhoneService | MultipleLines | InternetService | OnlineSecurity | OnlineBackup | DeviceProtection |
|---|--------|---------------|---------|------------|--------|--------------|------------------|-----------------|----------------|--------------|------------------|
| 0 | Female | 0 | Yes | No | 1 | No | No phone service | DSL | No | Yes | No |
| 1 | Male | 0 | No | No | 34 | Yes | No | DSL | Yes | No | Yes |
| 2 | Male | 0 | No | No | 2 | Yes | No | DSL | Yes | Yes | No |
| 3 | Male | 0 | No | No | 45 | No | No phone service | DSL | Yes | No | Yes |
| 4 | Female | 0 | No | No | 2 | Yes | No | Fiber optic | No | No | No |

```
In [ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 20 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 7043 non-null  object
1   SeniorCitizen          7043 non-null  int64
2   Partner                7043 non-null  object
3   Dependents             7043 non-null  object
4   tenure                 7043 non-null  int64
5   PhoneService           7043 non-null  object
6   MultipleLines          7043 non-null  object
7   InternetService        7043 non-null  object
8   OnlineSecurity         7043 non-null  object
9   OnlineBackup           7043 non-null  object
10  DeviceProtection       7043 non-null  object
11  TechSupport            7043 non-null  object
12  StreamingTV            7043 non-null  object
13  StreamingMovies        7043 non-null  object
14  Contract               7043 non-null  object
15  PaperlessBilling       7043 non-null  object
16  PaymentMethod          7043 non-null  object
17  MonthlyCharges         7043 non-null  float64
18  TotalCharges           7043 non-null  object
19  Churn                  7043 non-null  object
dtypes: float64(1), int64(2), object(17)
memory usage: 1.1+ MB
```

```
In [ ]: df = df[df["TotalCharges"] != " "]
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'])
```

Data Pre-Processing

```
In [ ]: import pandas as pd

# Load your DataFrame `df` here

# Target Transformation
df['Churn'] = df['Churn'].map({"No": 0, "Yes": 1})
```

```

# Min-Max Scaling for 'TotalCharges'
df = df[df["TotalCharges"] != " "]
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'])
df['TotalCharges'] = (df['TotalCharges'] - df['TotalCharges'].min()) / (df['TotalCharges'].max() - df['TotalCharges'].min())

# Min-Max Scaling for 'MonthlyCharges'
df = df[df["MonthlyCharges"] != " "]
df['MonthlyCharges'] = pd.to_numeric(df['MonthlyCharges'])
df['MonthlyCharges'] = (df['MonthlyCharges'] - df['MonthlyCharges'].min()) / (df['MonthlyCharges'].max() - df['MonthlyCharges'].min())

# Min-Max Scaling for 'tenure'
df = df[df["tenure"] != " "]
df['tenure'] = pd.to_numeric(df['tenure'])
df['tenure'] = (df['tenure'] - df['tenure'].min()) / (df['tenure'].max() - df['tenure'].min())

# One-Hot Encoding for categorical columns
df = pd.get_dummies(df)

data = df
data.head()

```

Out[]:

| | SeniorCitizen | tenure | MonthlyCharges | TotalCharges | Churn | gender_Female | gender_Male | Partner_No | Partner_Yes | Dependents_No | ... | Stream |
|---|---------------|----------|----------------|--------------|-------|---------------|-------------|------------|-------------|---------------|-----|--------|
| 0 | 0 | 0.000000 | 0.115423 | 0.001275 | 0 | 1 | 0 | 0 | 1 | 1 | ... | |
| 1 | 0 | 0.464789 | 0.385075 | 0.215867 | 0 | 0 | 1 | 1 | 0 | 1 | ... | |
| 2 | 0 | 0.014085 | 0.354229 | 0.010310 | 1 | 0 | 1 | 1 | 0 | 1 | ... | |
| 3 | 0 | 0.619718 | 0.239303 | 0.210241 | 0 | 0 | 1 | 1 | 0 | 1 | ... | |
| 4 | 0 | 0.014085 | 0.521891 | 0.015330 | 1 | 1 | 0 | 1 | 0 | 1 | ... | |

5 rows × 46 columns

Model Training

```
In [ ]: x_columns = data.columns.drop('Churn').tolist()
x = data[x_columns]
y = data['Churn']
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.3, random_state=42, shuffle=True)
```

SVM

```
In [ ]: svcModel = SVC()
svcModel.fit(x_train, y_train)
accuracy = svcModel.score(x_test, y_test)
print("Accuracy:", accuracy)
```

Accuracy: 0.7900473933649289

Decision Tree

```
In [ ]: dtModel = DecisionTreeClassifier()
dtModel.fit(x_train, y_train)
accuracy_dt = dtModel.score(x_test, y_test)
print("Decision Tree Accuracy:", accuracy_dt)
```

Decision Tree Accuracy: 0.7327014218009479

Random Forest

```
In [ ]: rfModel = RandomForestClassifier()
rfModel.fit(x_train, y_train)
accuracy_rf = rfModel.score(x_test, y_test)
print("Random Forest Accuracy:", accuracy_rf)
```

Random Forest Accuracy: 0.7777251184834123

Adaboost

```
In [ ]: adaModel = AdaBoostClassifier()
adaModel.fit(x_train, y_train)
accuracy_ada = adaModel.score(x_test, y_test)
print("AdaBoost Accuracy:", accuracy_ada)
```

AdaBoost Accuracy: 0.7909952606635071

HyperParameter Tuning

Random Forest Classifier

GridSearchCV

```
In [ ]: param_grid_rf = {'n_estimators': [100, 200, 300], 'criterion': ['gini', 'entropy'],
                        'max_depth': [None, 10, 20], 'min_samples_split': [2, 5, 10],
                        'min_samples_leaf': [1, 2, 4]}
grid_search_rf = GridSearchCV(RandomForestClassifier(), param_grid_rf, cv=5)
grid_search_rf.fit(x_train, y_train)
best_params_rf = grid_search_rf.best_params_
print("Best Parameters for Random Forest:", best_params_rf)
```

Best Parameters for Random Forest: {'criterion': 'entropy', 'max_depth': 20, 'min_samples_leaf': 4, 'min_samples_split': 10, 'n_estimators': 300}

RandomizedSearch CV

```
In [ ]: param_dist_ada = {'n_estimators': [50, 100, 150], 'learning_rate': [0.01, 0.1, 1.0]}
random_search_ada = RandomizedSearchCV(AdaBoostClassifier(), param_dist_ada, cv=5, n_iter=10)
random_search_ada.fit(x_train, y_train)
best_params_rand_ada = random_search_ada.best_params_
print("Best Parameters for AdaBoost (RandomizedSearchCV):", best_params_rand_ada)
```

c:\Python311\Lib\site-packages\sklearn\model_selection_search.py:307: UserWarning: The total space of parameters 9 is smaller than n_iter=10. Running 9 iterations. For exhaustive searches, use GridSearchCV.

warnings.warn(

Best Parameters for AdaBoost (RandomizedSearchCV): {'n_estimators': 150, 'learning_rate': 0.1}

Decision Tree classifier

GridSearch CV

```
In [ ]: param_grid_dt = {'criterion': ['gini', 'entropy'], 'splitter': ['best', 'random'],
                        'max_depth': [None, 10, 20], 'min_samples_split': [2, 5, 10],
                        'min_samples_leaf': [1, 2, 4]}
grid_search_dt = GridSearchCV(DecisionTreeClassifier(), param_grid_dt, cv=5)
grid_search_dt.fit(x_train, y_train)
best_params_dt = grid_search_dt.best_params_
print("Best Parameters for Decision Tree:", best_params_dt)
```

Best Parameters for Decision Tree: {'criterion': 'gini', 'max_depth': 10, 'min_samples_leaf': 4, 'min_samples_split': 2, 'splitter': 'random'}

Randomized Search CV

```
In [ ]: param_dist_dt = {'criterion': ['gini', 'entropy'], 'splitter': ['best', 'random'],
                        'max_depth': [None, 10, 20], 'min_samples_split': [2, 5, 10],
                        'min_samples_leaf': [1, 2, 4]}
random_search_dt = RandomizedSearchCV(DecisionTreeClassifier(), param_dist_dt, cv=5, n_iter=10)
random_search_dt.fit(x_train, y_train)
best_params_rand_dt = random_search_dt.best_params_
print("Best Parameters for Decision Tree (RandomizedSearchCV):", best_params_rand_dt)
```

Best Parameters for Decision Tree (RandomizedSearchCV): {'splitter': 'random', 'min_samples_split': 5, 'min_samples_leaf': 4, 'max_depth': 10, 'criterion': 'gini'}

SVM Classifier

GridSearch CV

```
In [ ]: param_grid_svc = {'C': [0.1, 1, 10], 'kernel': ['linear', 'rbf'], 'gamma': ['scale', 'auto']}
grid_search_svc = GridSearchCV(SVC(), param_grid_svc, cv=5)
grid_search_svc.fit(x_train, y_train)
```

```
best_params_svc = grid_search_svc.best_params_  
print("Best Parameters for SVC:", best_params_svc)
```

Best Parameters for SVC: {'C': 1, 'gamma': 'auto', 'kernel': 'rbf'}

Randomized Search CV

```
In [ ]: from scipy.stats import uniform  
param_dist_svc = {'C': uniform(loc=0, scale=10), 'kernel': ['linear', 'rbf'], 'gamma': ['scale', 'auto']}  
random_search_svc = RandomizedSearchCV(SVC(), param_dist_svc, cv=5, n_iter=10)  
random_search_svc.fit(x_train, y_train)  
best_params_rand_svc = random_search_svc.best_params_  
print("Best Parameters for SVC (RandomizedSearchCV):", best_params_rand_svc)
```

Best Parameters for SVC (RandomizedSearchCV): {'C': 4.075240301647072, 'gamma': 'auto', 'kernel': 'rbf'}